

A large, faint, circular watermark of the Stanford University seal is centered in the background. The seal features a redwood tree in the center, surrounded by the text "STANFORD JUNIOR UNIVERSITY" at the top and "DIE LUFT DER FREIHEIT WEHT" at the bottom. The year "1891" is at the bottom of the seal. The seal is bordered by a diamond pattern.

EE382A – Spring 2025


**Chapter 2:
Kestrel's Instruction Set
Architecture (1 of 5)**

Kestrel's ISA, first part

- Overview of Kestrel's ISA
- First program: compiling, running, debugging
- Second program: PE numbering, loops
- Third program: move, I/O
- Kestrel's streaming I/O system
- Programming assignment 1: PE count, PE alternate

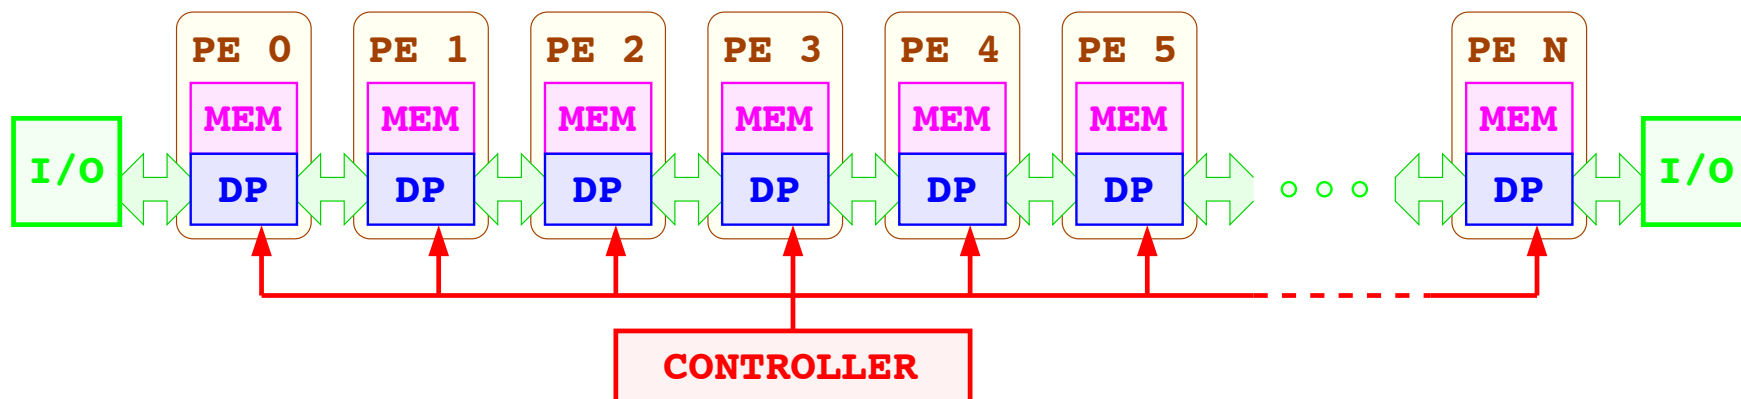


UCSC Kestrel

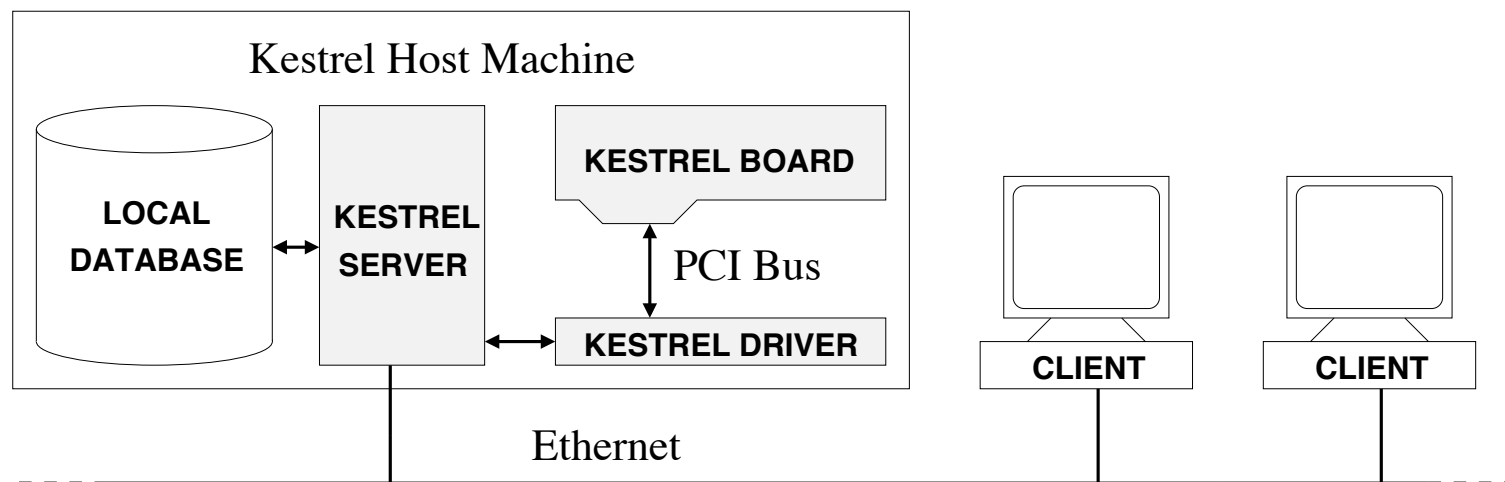


- 512 8-bit PEs, 256 bytes per PE, 20 MHz
- Linear SIMD array with wrap-around
- About 25 8-bit Gops peak (MAA = 2.5 ops)

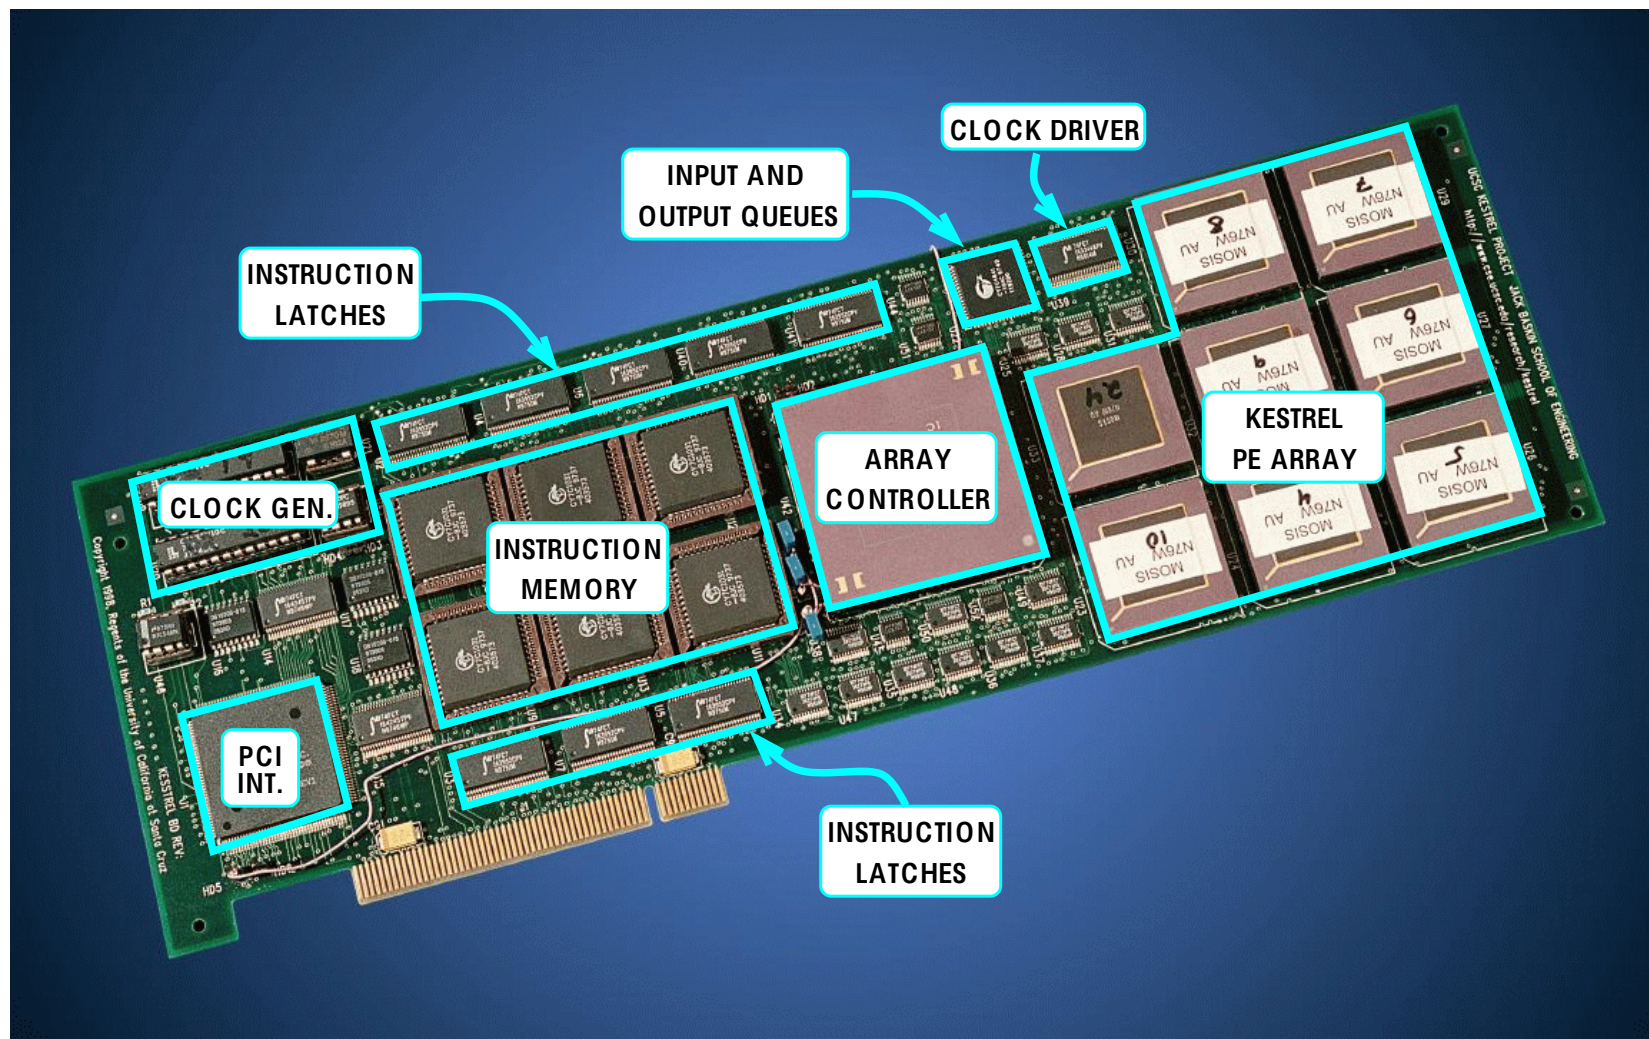
UCSC Kestrel: Linear SIMD array



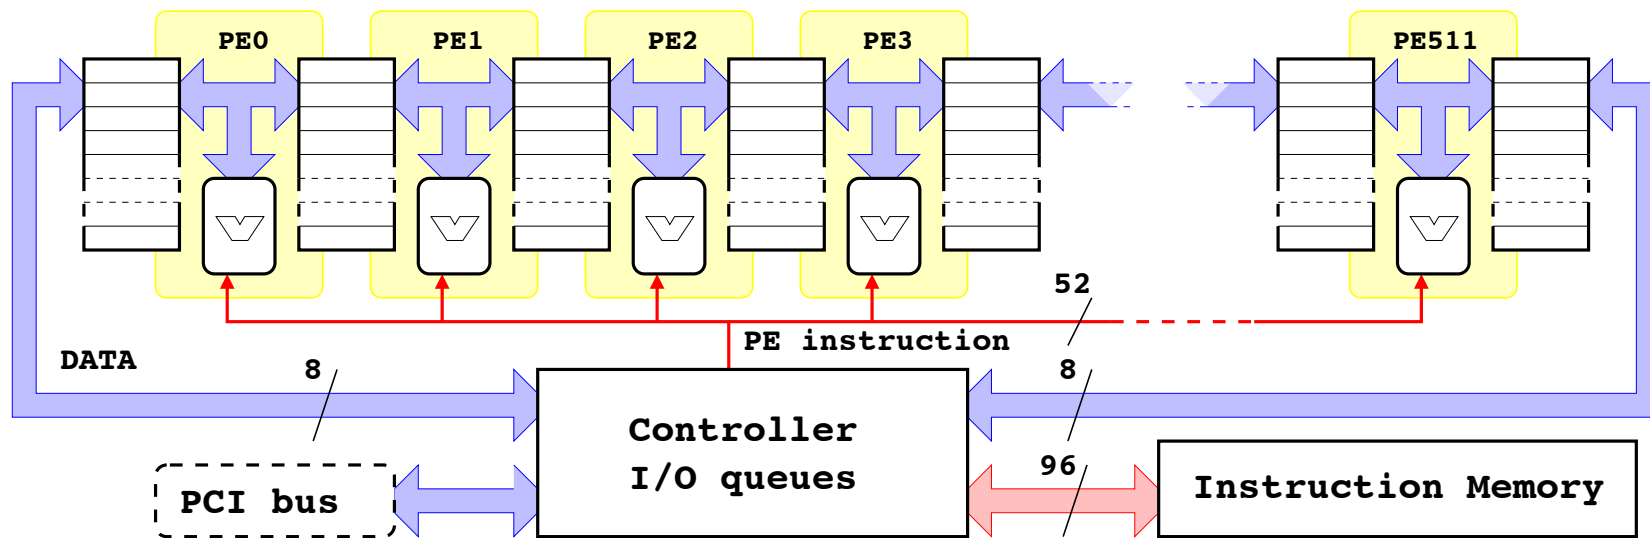
UCSC Kestrel: host and clients



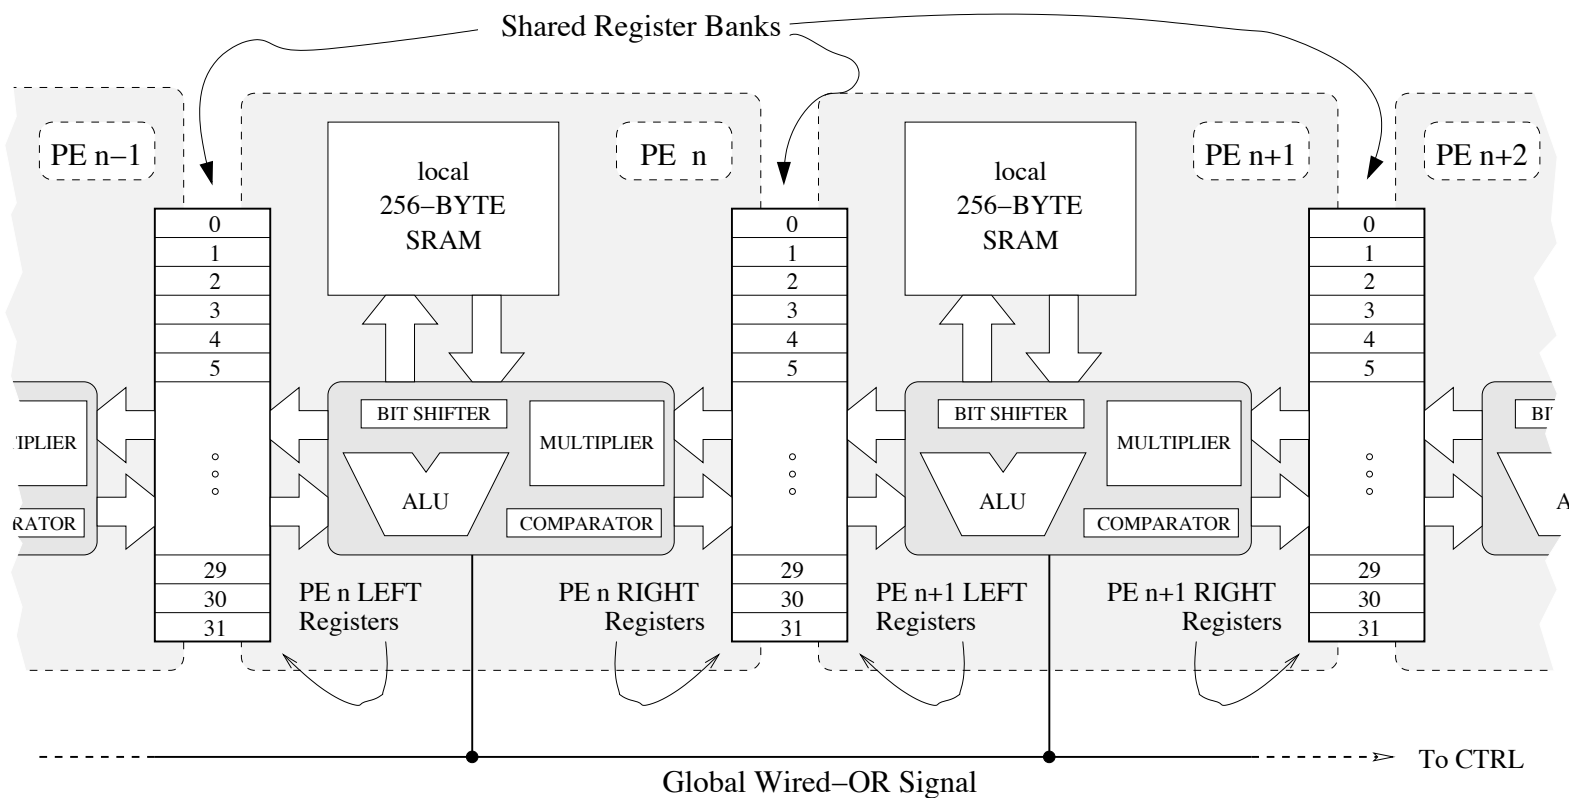
UCSC Kestrel: the board



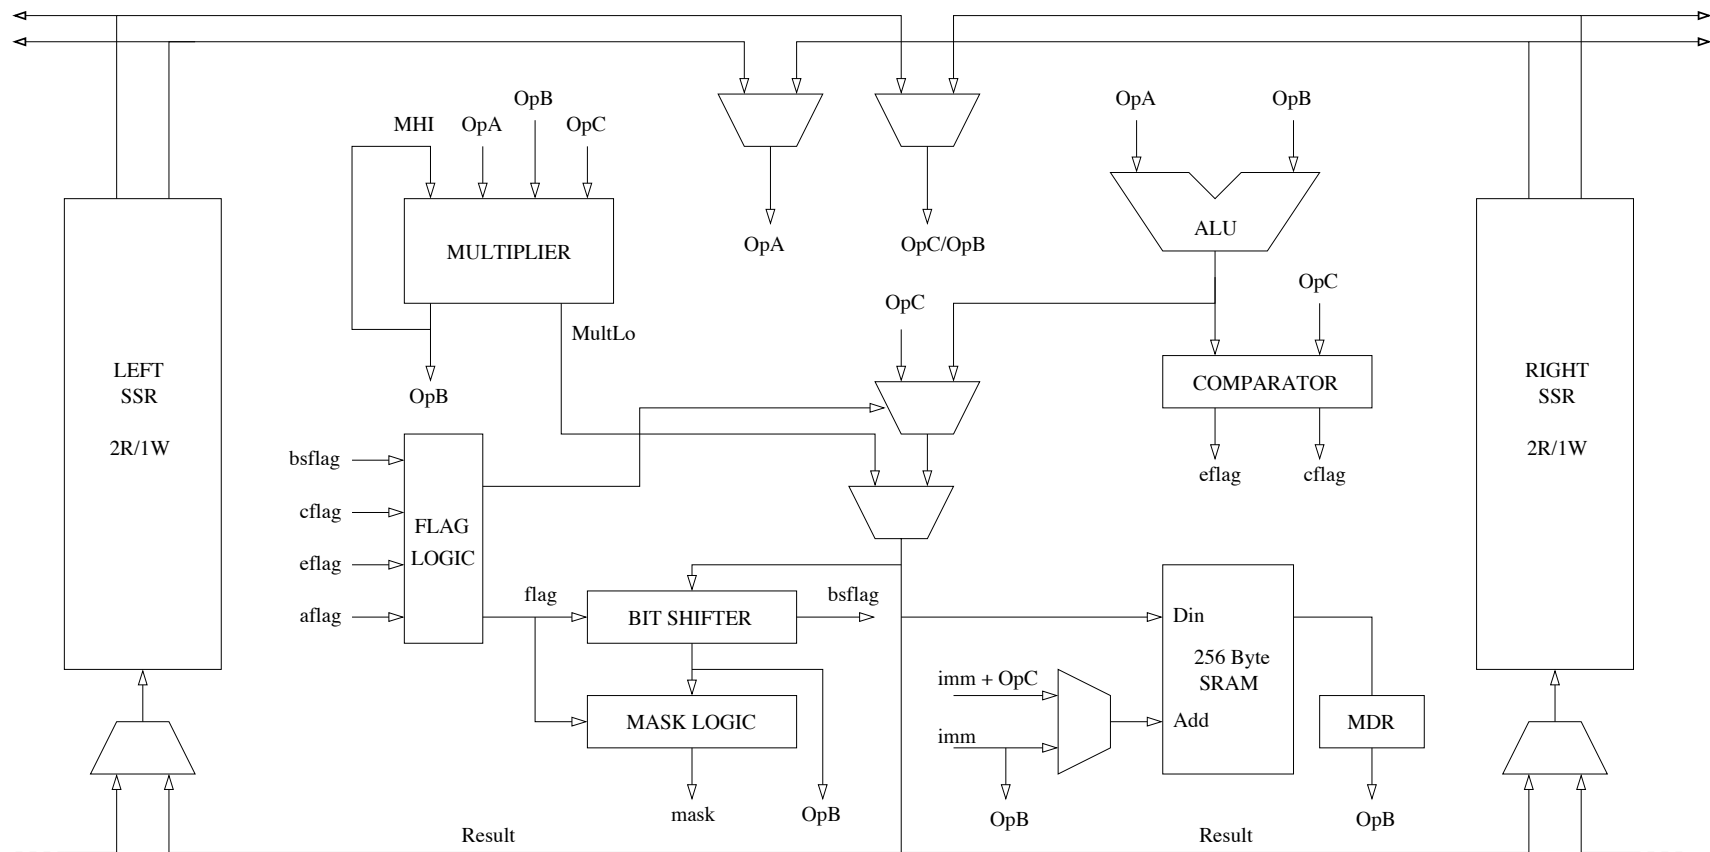
Kestrel's array structure



Kestrel array detail



Kestrel processing element detail



Kestrel ISA

- 8-bit word size, support for multi-precision operations
- Two *systolic shared register* banks per PE, 32 8-bit registers each
- 0-, 1-, 2-, 3-, 4-, or 5-operand instructions - high instruction-level parallelism with multiple functional units
- Load/store architecture
- 256 bytes of local, locally-addressable memory per PE
- Single data addressing mode
- Single-cycle implementation
- No floating-point unit or support
- Official manual: Richard Hughey, *KASM Assembly Manual*, Tech. Rep. UCSC-CRL-98-11, UC Santa Cruz, Sept. 1998 (a copy is provided in our course's website)



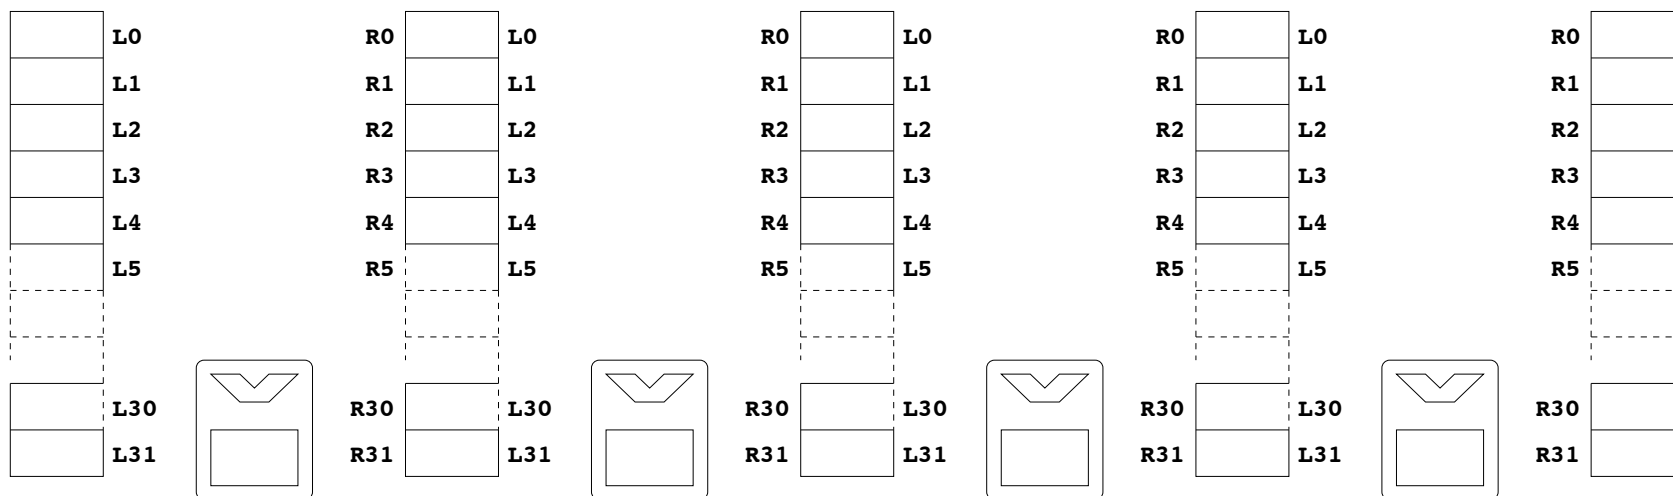
First Kestrel program

```
;*****  
; Program: kex01.kasm  
; First Kestrel program example  
; Andrea Di Blas, May 2005  
;*****  
start:      ; program execution starts here  
addzz      L0  
add         L0, L0, #3  
;*****
```

What does this program do?

What does `kex01.kasm` do?

Assume to run it on a 4-PE array like this:



```
start:    ; program execution starts here
addzz    L0
add      L0, L0, #3
```

Assembling the program

To install the Kestrel environment:

- Download the tar files **kestrelClient.tar.gz**, **kestrel_win32.tar.gz**, and **newkasm.tar.gz** from our web page.
- Follow the instructions to **make** the assembler, **newkasm**, and the client/simulator/debugger, **kestrel** (look at the README files). Make sure to **make clean** before **make**.

To assemble the program:

```
$ newkasm kex01.kasm
```

The output is an object (executable) called **kex01.ko**



Running the program

Simulating an array with 4 PEs:

```
$ kestrel kex01.ko kin kout 4
```

Where **kin** is the input file name and **kout** is the output file name. Both need to be specified even when not actually used (use any dummy input file for **kex01**).

If you forget the syntax:

```
$ kestrel -h
```

Except that the board option is not available (simulator only)



Addition instructions

ADD writes in **Rdst** the sum of two operands

ADD Rdst, OpB, OpA

ADDXX writes in **Rdst** sum of an operand with itself

ADDXX Rdst, [OpA or OpB]

ADDXZ writes in **Rdst** sum of an operand with zero

ADDXZ Rdst, [OpA or OpB]

ADDZZ writes in **Rdst** sum of zero with zero

ADDZZ Rdst

Operand A, operand B, operand C

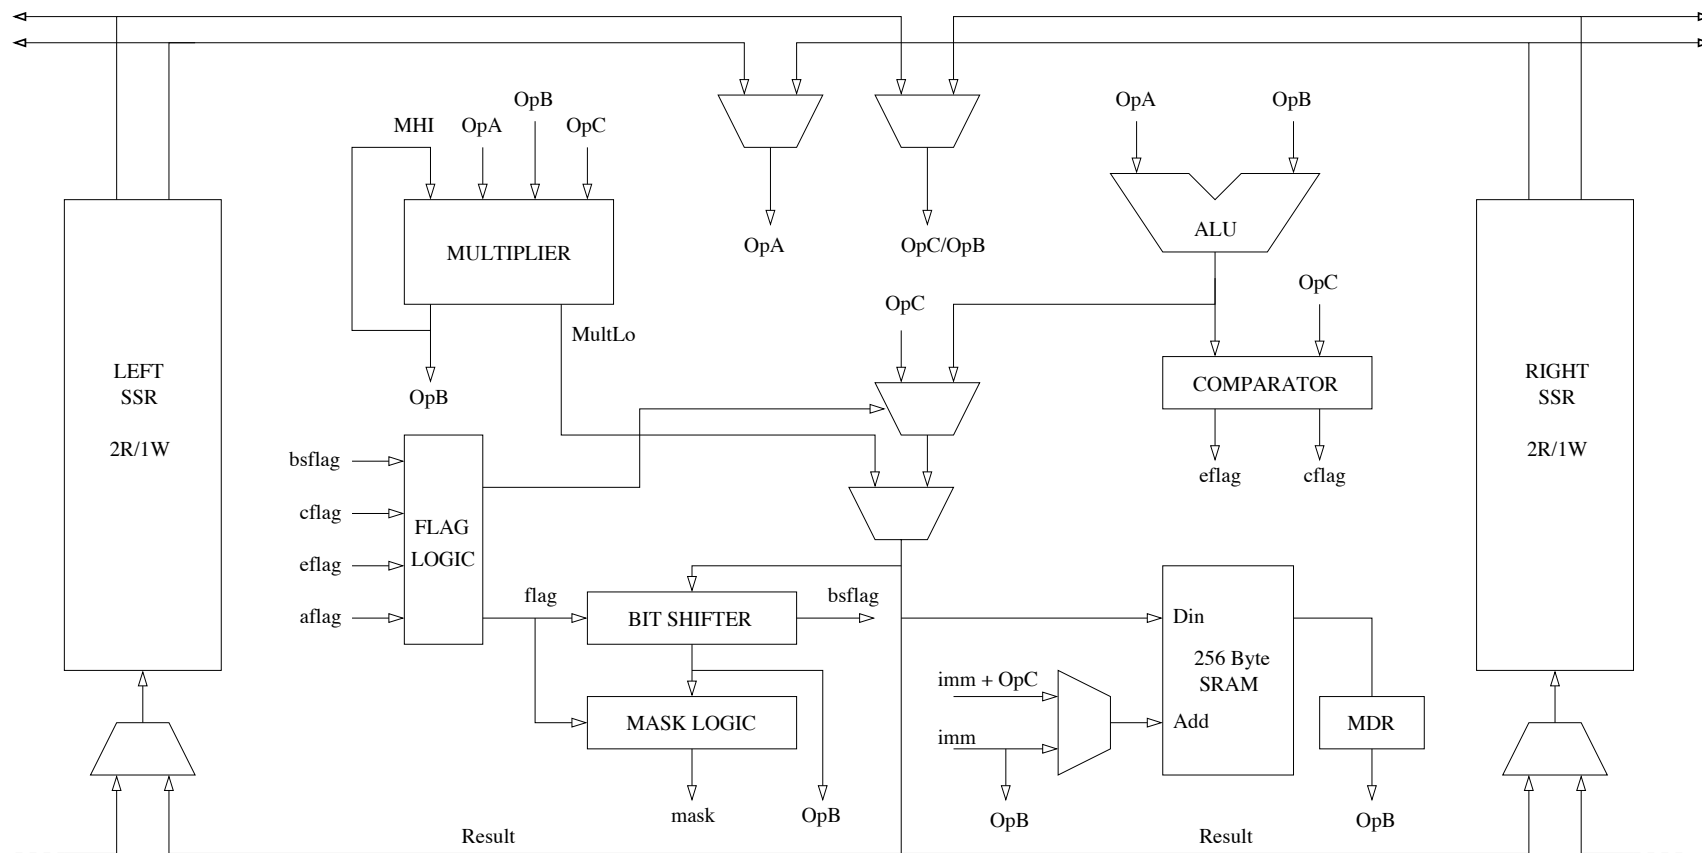
Three operand busses, A, B, and C (fig. on page 2.5).

Operand A and operand C can only be registers.

Operand B can assume one of the following values:

- Operand C (e.g. **L17**)
- Sign extension of operand C (e.g. **sL17**)
- Memory Data Register (**mdr**)
- Sign extension of **mdr** (**smdr**)
- Multiplier high byte (**mhi**)
- Sign extension of **mhi** (**smhi**)
- Bit shifter (**bs**)
- Instruction immediate (e.g. **#73**)

Kestrel processing element detail again



Running kex01 in the debugger

Start the debugger, simulate 4 PEs:

```
$ kestrel -debug kex01.ko kin kout 4
```

```
Kestrel Run Time Environment (compiled 05/08/99_17:57:43)
```

```
Copyright (c) 1998 Regents of the University of California
```

```
kestrel rte: Starting the Kestrel Serial Simulator.
```

```
kestrel rte: kex01.ko: end of program detected at instruction 3
```

```
kestrel rte: Starting the Kestrel Debugger (kdb)
```

(continued on the next page)

The debugger main menu

@Main Menu:

run	- Run your program
step	- single Step through your program
examine	- Examine a specific PE
range	- examine single value across a Range of PEs
controller	- examine state of the controller
list	- List most recent instructions
breakpoint	- set, change or remove Breakpoints
format	- change the display format
precision	- change the Precision
settings	- print current format information
history	- print history buffer of commands
menu	- display this Menu
quit	- Quit kdb
dump	- Single step program, dumping state to file 'state.dump'

@Main Menu:

kdb>



Stepping through kex01: the range menu

```
kdb> range
```

```
@Main Menu > Range Menu:
```

```
reg <#> [- <#>] - examine a range of registers in each SSR
sram <#> [- <#>] - examine a range of sram locations in each PE
masklatch        - view the maskLatch
minlatch         - view the minLatch
bslatch          - view the bsLatch
mdrlatch         - view the mdrLatch
clatch           - view the cLatch
multhilatch      - view the multhilatch
eqlatch          - view the eqLatch
menu             - display this Menu
back            - back to Main menu
quit            - Quit kdb
```

```
kdb>
```

NOTE that **range** is reliable, while **examine** reports incorrect values - do not use.

Looking at register 0

```
kdb> reg 0
```

```
kestrel rte: get register state: end of program detected at instruction 455
```

```
kestrel rte: running program get register state at 6
```

```
kestrel rte: end of program reached at 455
```

```
QIN DATA: 1 QOUT DATA: 256
```

```
reg    0 values from 0 to 3 are:
```

```
0:    0,    0    1:    0,    0    2:    0,    0    3:    0,    0
```

```
kdb>
```

Stepping through kex01

```
kdb> b
```

```
@Main Menu:
```

```
kdb> step
```

```
kestrel rte: Running kex01.ko at 0
```

```
kestrel rte: Status Register: kestrel rte: program stepped, now at address 0
```

```
kdb> step
```

```
kestrel rte: Status Register: kestrel rte: program stepped, now at address 1
```

```
kdb> step
```

```
kestrel rte: Status Register: kestrel rte: program stepped, now at address 2
```

```
kdb>
```

Here's the code again:

```
start:                                ; address 0
addzz      L0                        ; address 1
add         L0, L0, #3                ; address 2
[Program Termination]                ; address 3
```



Stepping through kex01

```
kdb> reg 0
```

```
QIN DATA: 11 QOUT DATA: 0
```

```
kestrel rte: get register state: end of program detected at instruction 455
```

```
kestrel rte: running program get register state at 6
```

```
kestrel rte: end of program reached at 455
```

```
QIN DATA: 1 QOUT DATA: 256
```

```
reg    0 values from 0 to 3 are:
```

```
0:    0,    0  1:    0,    0  2:    0,    0  3:    0,    0
```

```
kdb>
```


Stepping through kex01: run to end

```
kdb> b
```

```
@Main Menu:
```

```
kdb> step
```

```
kestrel rte: Status Register: kestrel rte: program stepped, now at address 2
```

```
kdb> run
```

```
kestrel rte: continuing execution of kex01.ko at 2.
```

```
kestrel rte: end of program reached at 3
```

```
QIN DATA: 3 QOUT DATA: 0
```

```
kestrel rte: program kex01.ko completed at 3 with 0 output bytes.
```

```
kestrel rte: program used only 0 of 2 input data bytes.
```

```
kdb> range
```

```
@Main Menu > Range Menu:
```

[The Range Menu]

Looking at final content of register 0

```
kdb> reg 0
```

```
kestrel rte: get register state: end of program detected at instruction 455
```

```
kestrel rte: running program get register state at 6
```

```
kestrel rte: end of program reached at 455
```

```
QIN DATA: 1 QOUT DATA: 256
```

```
reg    0 values from 0 to 3 are:
```

```
0:    3,    3    1:    3,    3    2:    3,    3    3:    3,    0
```

```
kdb>
```

What does `kex01.ko` look like?

```
$ cat kex01.ko
```

```
b01000000000000200000218001
```

```
bc900000000000200000018000
```

```
e09001c0000003200000018000
```

```
b01000000000000200000118003
```

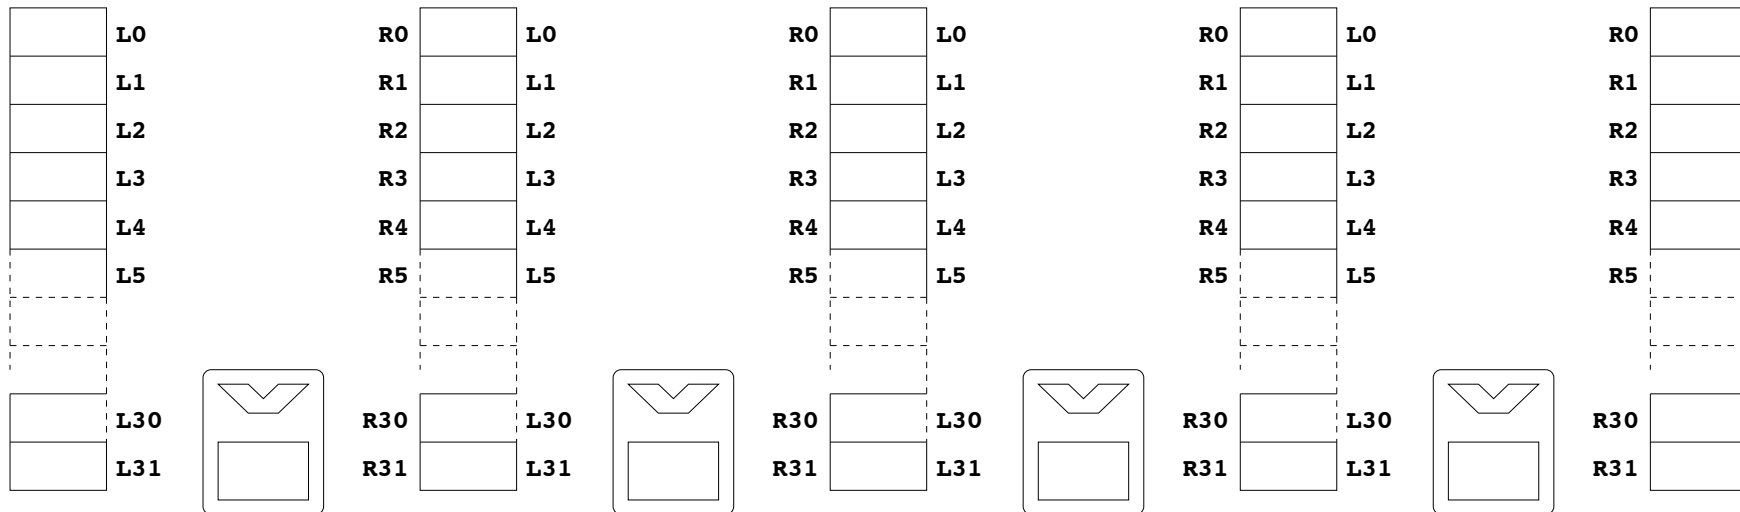
```
$
```

Second Kestrel program

```
;*****  
; Program: kex02.kasm  
; Kestrel program example  
;*****  
start:    ; program execution starts here  
addxz    L0, #1  
add      R0, L0, #1  
add      R0, L0, #1  
add      R0, L0, #1  
;*****
```

What does this program do?

Tracing kex02



addxz L0, #1

add R0, L0, #1

add R0, L0, #1

add R0, L0, #1

Loops

The controller has a loop counter stack that supports up to 15 nested loops.

BEGINLOOP pushes a new loop counter onto the stack

beginloop CImm

Where **CImm** is a 16-bit immediate value, unsigned.

ENDLOOP decrements the current loop counter (TOS) and branches back to the instruction following the corresponding **beginloop** if the count is greater than zero. If the count is zero, pops it off the stack and proceeds with the instruction following the **endloop**.

endloop

NOTE: the loop counter is NOT available as a variable. Also, the branch target address for **endloop** is encoded into the controller's 16-bit immediate field as an *absolute* reference to the instruction number, so no loop can *begin* at addresses higher than 64K.



Example of loop

```
;*****  
; Program: kex03.kasm  
; Kestrel program example  
; Same as kex02 but with a loop  
; Andrea Di Blas, May 2005  
;*****  
start:    ; program execution starts here  
addxz    L0, #1  
beginloop 3  
        add    R0, L0, #1  
endloop  
;*****
```



Controller and array immediate

Two different immediate fields:

Controller immediate:

- 16 bit
- unsigned
- syntax: **12345** (only used with **beginloop**)

Array immediate

- 8 bits
- unsigned or signed (two's complement)
- syntax: **#123**




The MOVE instruction

MOVE move an operand to a destination:


MOVE Dst, OpA or OpB

Example: Source and destination are registers on same side:



The MOVE instruction: another example

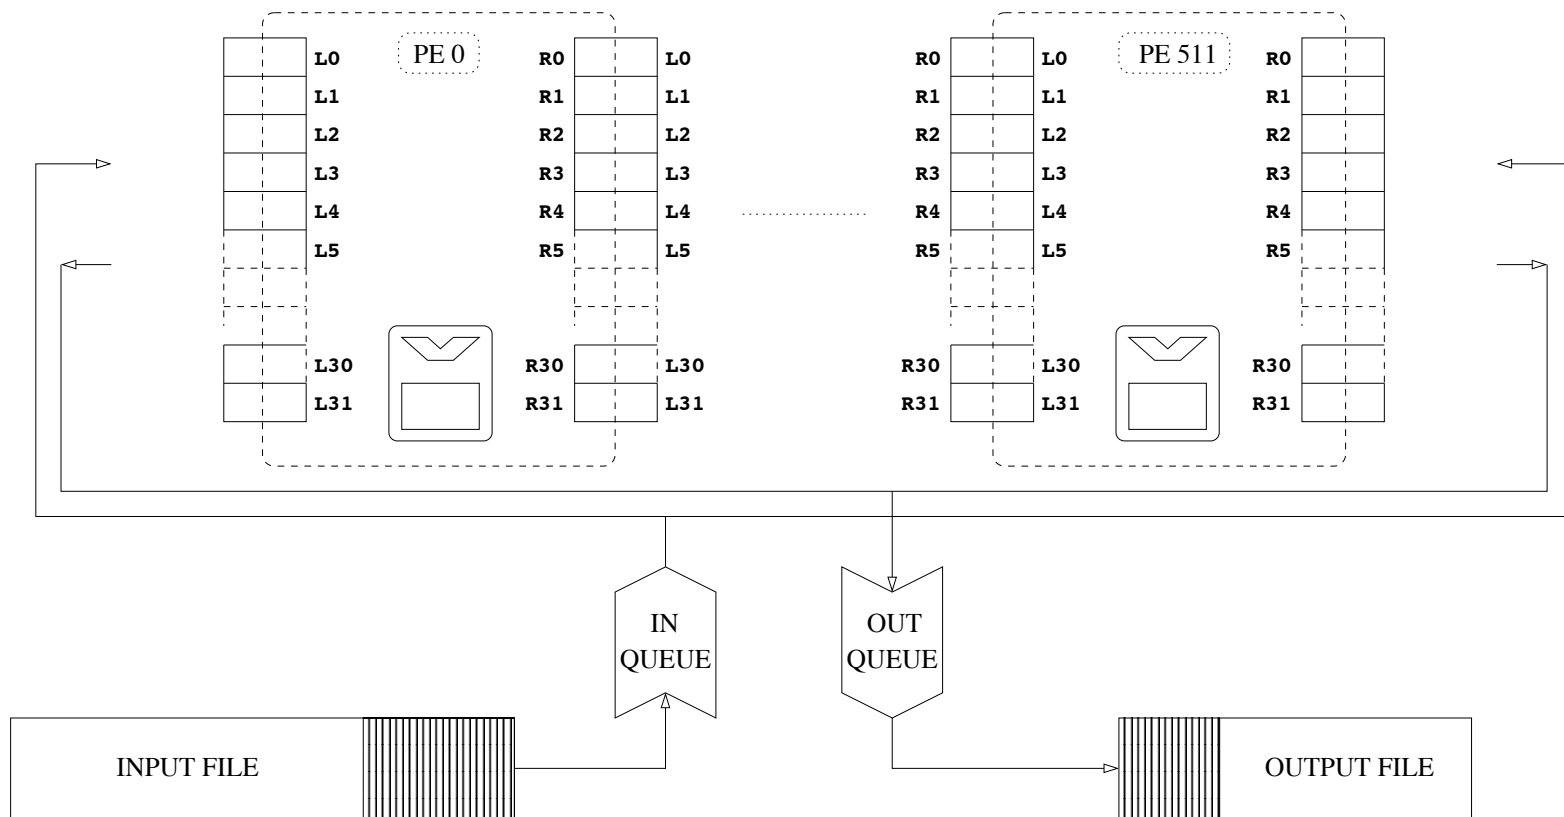
Source and destination are registers on different sides:



What happens at the edges?

Kestrel I/O system

There are always two files: an input file and an output file.



The input file is a FIFO called Input Queue, the output file is a FIFO called Output Queue

File input: **qtoarr**

QTOARR “Queue **TO** **ARR**ay” writes the next byte in the input file into the destination register

<some instr> Rdst, ... qtoarr

- If **Rdst** is a *left* register, the value will be written into the right register bank of the rightmost PE. If **Rdst** is a *right* register, the value will be written into the left register bank of the leftmost PE.
- The registers are written regardless of the *mask* of the PE being written (explained below).
- Input and output can overlap, so **qtoarr** and **arrtoq** can appear in the same instruction.

File output: **arrtoq**

ARRTOQ “**ARR**ay **TO** **Q**ueue” writes the result of the instruction to the output file as well as into the destination register and into the controller’s *scratch* register.

<some instr> Rdst, ... arrtoq

- If **Rdst** is a *left* register, the value comes from the leftmost PE in the array. If **Rdst** is a *right* register, the value comes from the rightmost PE in the array.
- If the outputting PE is ***masked***, no output is produced (explained later).

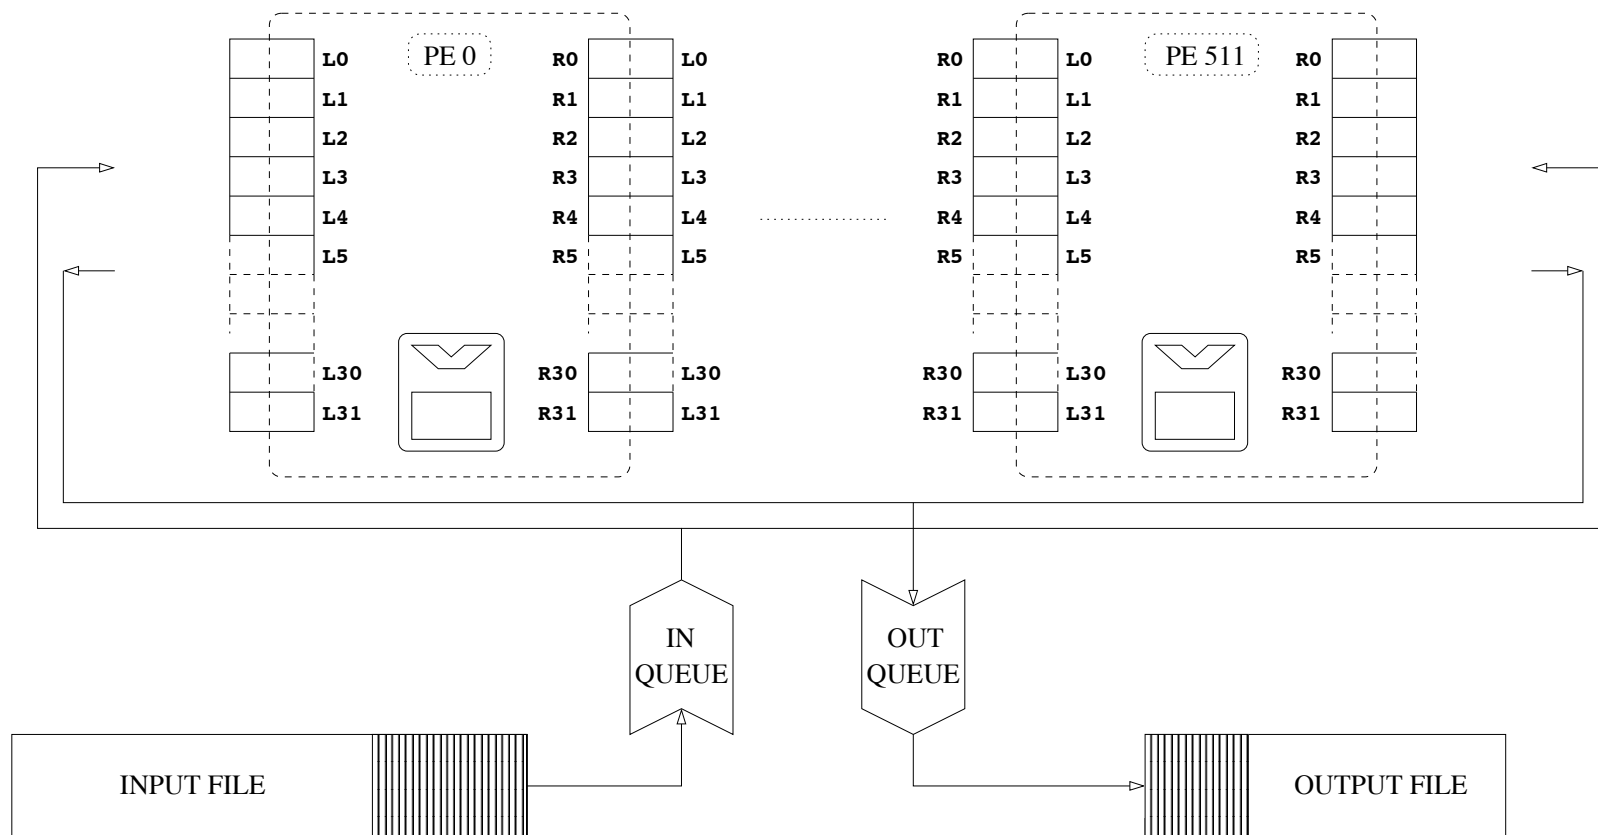


File input and output overlap

- Input and output can overlap, so **qtoarr** and **arrtoq** can appear in the same instruction.
- That is the maximum extent of I/O parallelism that we have here, no multiply opened files with parallel access




I/O example




1. **ADD R2, L0, R4, arrtoq**
2. **SUB R2, L0, R4, qtoarr**
3. **MOVE L5, R5, qtoarr, arrtoq**

Stream I/O details


MOVE R0, L0, qtoarr




MOVE L0, R0, qtoarr



MOVE R0, L0, arrtoq




MOVE L0, R0, arrtoq



Stream I/O: output “anomaly”


Input can be modeled as if there was an additional “shadow PE” at the end of the array where the data is input from.

MOVE RO, L0, qtoarr




Output, however, behaves differently and counter-intuitively: in this example, the rightmost RO is not written to to output file. The reason is in how inter-chip communication is implemented at the hardware level.

MOVE RO, L0, arrtoq




Stream I/O: more examples


MOVE L0, L0, qtoarr




MOVE R0, R0, qtoarr



MOVE L0, L0, arrtoq



MOVE R0, R0, arrtoq



Program with I/O

```

;*****
; Program: kex04.kasm
; Kestrel program example
;*****

start:
bsclearm
addzz      L0
beginloop  8
    move    R0, L0, qtoarr
endloop
add        R0, R0, #1
beginloop  8
    move    L0, R0, arrtoq
endloop
end:
;*****

```

What does this program do?



Input and output file format

The input and output file format can be specified individually for each file with the option:

- iformat <format>** for the input file, or
- oformat <format>** for the output file.

where **<format>** can be:

- **decimal** for ASCII decimal data, newline separated (default)
- **octal** for ASCII octal data, newline separated
- **hex** for ASCII hex data, newline separated
- **binary** for raw binary data



Kestrel I/O file format example

Decimal	Octal	Hex	Raw
33\n	41\n	21\n	!
34\n	42\n	22\n	''
45\n	55\n	2D\n	-
56\n	70\n	38\n	8
9\n	11\n	9\n	\t
78\n	116\n	4E\n	N
1\n	1\n	1\n	[1]
2\n	2\n	2\n	[2]
...

Programming assignment 1: PE count

Since when running a program with the simulator we can specify the number of PEs on the command line, how can a *program* figure out *at run time* how many PEs are available?

Write a program that stores in the register L0 of all PEs the number of PEs that are being simulated, assuming it is at most 255. You can only use instructions that we have seen so far in this chapter.

Make sure it works with all corner cases, such as 1 and 255.



PE count example

```
examples$ kestrel -debug kPA01_PEmount.ko kin kout 181
...
@Main Menu:kdb> run
...
kdb> range
...
kdb> reg 0
...
reg    0 values from 0 to 180 are:
0: 181, 181    1: 181, 181    2: 181, 181    3: 181, 181
...
```

Programming assignment 1: PE alternate

Write a program that reads 8 numbers from a file, one byte each, and then writes them to the output file alternatively first the first one read, then the last one read, then the second one read, then the second to last one read, and so on. Example:

Input file: 1 2 3 4 5 6 7 8

Output file: 1 8 2 7 3 6 4 5

You can choose the number of PEs that should be used, but you should use as many PE as possible to make your code as simple and efficient as possible (specify NPES in the README file included with your submission).



Code grading criteria

25% Code compiles and runs, and runs simple input files correctly.

25% Code runs corner cases correctly.

25% Code runs “secret” input files correctly.

25% General code *quality* and *clarity*.



Clean code

- Is **easy to read**, to understand, to debug.
- Has **meaningful variable names** and appropriate **comments**: explain the “why”, summarize the “what.”
- Is **modular**, but without being too modular. e.g. some meaningful macros are good, too many macros with cryptic names are bad.
- Is **well aligned**, has columns that don't exceed 80 characters (Oracle, Amazon, and Google C Coding Standard!)
- **Doesn't jump around** too much and has clear entry and exit points from functions/blocks.



Code efficiency

- **Main principle: code should use all available parallelism.**
- **HOWEVER**, in assembly in general and especially in Kestrel, often efficiency and clarity go against each other. If you use all possible optimizations offered by ILP, your code becomes easily unreadable. Therefore, optimizing all the way is the right approach only when it's worth it, usually in inner loops.
- For example, merging the PEID assignment with input file read does save instructions, but when you have to deal with much bigger pieces of code, saving 512 clock cycles may be less important than keeping these two basic operations clearly separate.

