

The background of the slide features a large, faint, circular watermark of the Stanford University seal. The seal contains a redwood tree in the center, surrounded by the text "STANFORD JUNIOR UNIVERSITY" at the top and "DIE LUFT DER FREIHEIT WEHT" at the bottom. The year "1891" is inscribed at the very bottom of the seal.

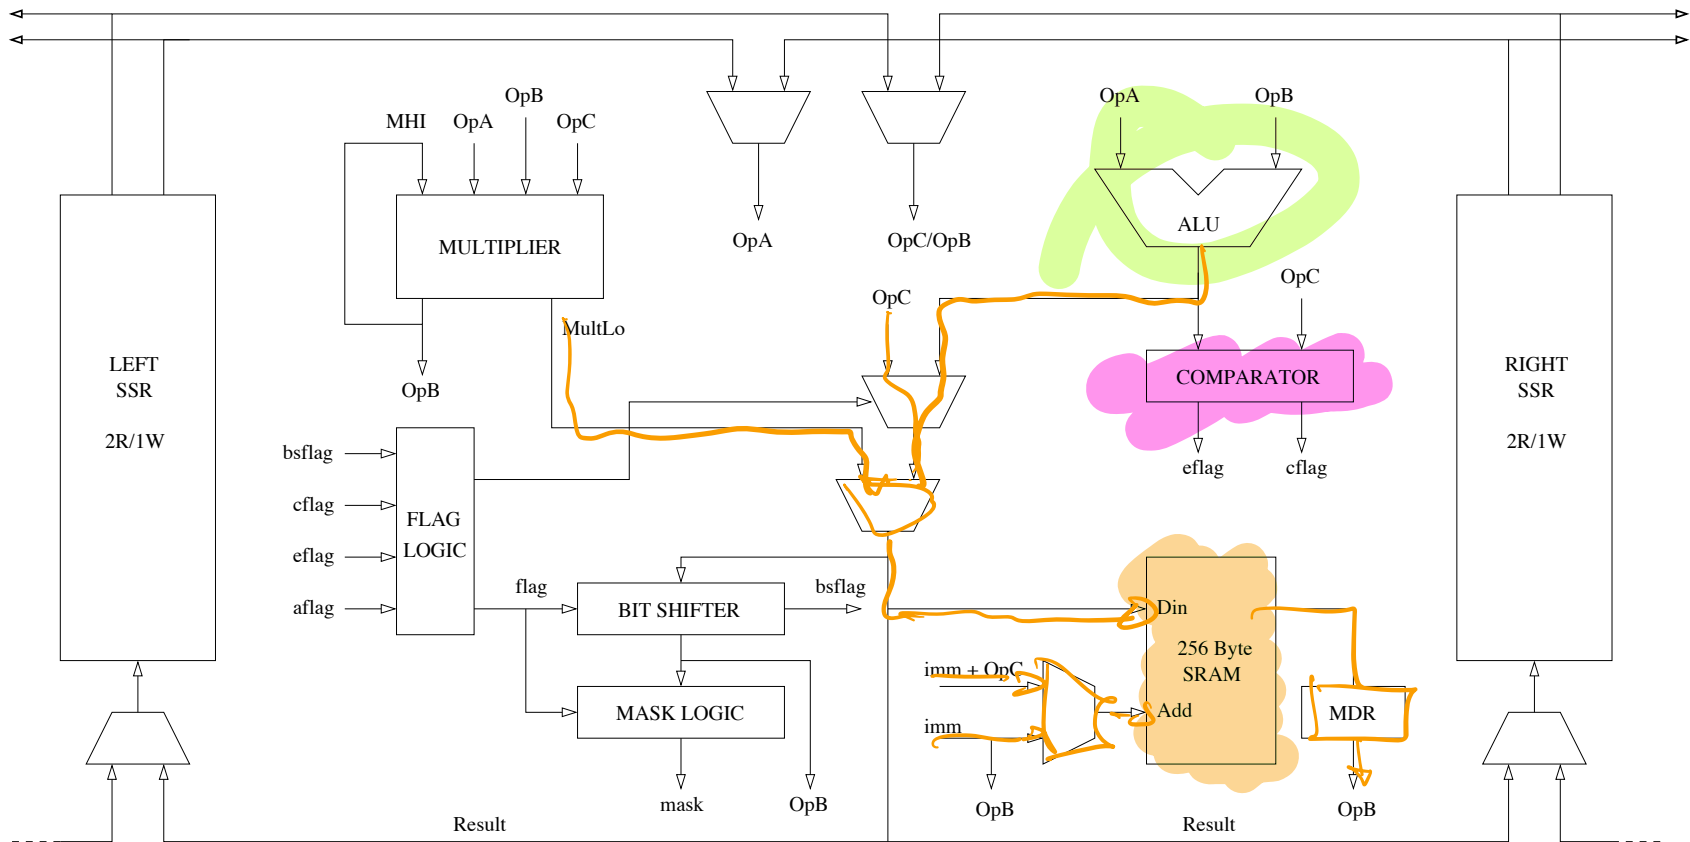
EE382A – Spring 2025

**Chapter 5:
Kestrel's Instruction Set
Architecture (3 of 5)**

Kestrel's ISA, third part

- Bitwise logical operations
- Multi-precision addition, subtraction, selection
- Multi-precision comparison
- Controller scratch register and run-time loop counter
- Local memory

Kestrel PE diagram



Bitwise logical operations

NOP No ALU operation

nop

INVERT Bitwise inversion (1's complement)

invert Rdst, {OpA | OpB}

AND, NAND, NOR, OR, XNOR, XOR Bitwise logical operations

{and | nand | nor | or | xnor | xor} Rdst, OpB, OpA

Subtraction

SUB Difference of two operands: $\text{OpB} - \text{OpA}$

sub Rdst, OpB, OpA

SUBXZ Subtract zero from operand: $\{\text{OpA} \mid \text{OpB}\} - 0$

subxz Rdst, {OpA | OpB}

SUBZX Subtract operand from zero: $0 - \{\text{OpA} \mid \text{OpB}\}$

subzx Rdst, {OpA | OpB}


(needs to be fixed - it works like a **SUBZZ** now)

SUBZZ Subtract zero from zero: $0 - 0$

subzz Rdst

Multi-precision addition

ALU diagram:



Multi-precision addition

Use modifiers:

- **c0** Set the ALU carry-in input to zero (default).
- **c1** Set the ALU carry-in input to one.
- **mp** Set the ALU carry-in input to the carry-out of the previous operation.

Examples:

addzz c0 L3 ; sets L3 = 0 $L3 = 0 + 0 + 0$

addzz c1 L3 ; sets L3 = 1 $L3 = 0 + 0 + 1$


addxz c1 L7, L7 ; increments L7 by one $L7 = L7 + 1$

Multi-precision addition example

Assume $[L1, L0] = A$, $[R1, R0] = B$, produce $[L11, L10] = A + B$:

```
add      L10, L0, R0 ; Low byte
add      mp L11, L1, R1 ; High byte
addzz    mp L12      ; Overflow?
```

Difference between signed and unsigned: $A + B$ where $A = 20$, $B = -4$
Overflow detection?



Addition: overflow

Unsigned: overflow = carry out (**fbaco**)

Signed: overflow = ALU true sign \neq result sign (**fbats**)

Example on 8 bits:

UNSIGNED

$$\begin{array}{r}
 255 \quad 11111110 \\
 + \quad 1 \quad 00000000 \\
 \hline
 00000000
 \end{array}$$

SIGNED

$$\begin{array}{r}
 127 \quad 01111110 \\
 + \quad 1 \quad 00000000 \\
 \hline
 01000000 = -128
 \end{array}$$

No overflow when numbers have different signs.

Addition: overflow

From RISC-V Specification, User mode:

We did not include special instruction-set support for overflow checks on integer arithmetic operations in the base instruction set, as many overflow checks can be cheaply implemented using 18 Volume I: RISC-V Unprivileged ISA V20191213 RISC-V branches. Overflow checking for unsigned addition requires only a single additional branch instruction after the addition: **add t0, t1, t2; bltu t0, t1, overflow**. For signed addition, if one operand's sign is known, overflow checking requires only a single branch after the addition: **addi t0, t1, +imm; blt t0, t1, overflow**. This covers the common case of addition with an immediate operand. For general signed addition, three additional instructions after the addition are required, leveraging the observation that the sum should be less than one of the operands if and only if the other operand is negative. **add t0, t1, t2 slti t3, t2, 0 slt t4, t0, t1 bne t3, t4, overflow**



Multi-precision subtraction

Example: 10 - 30 on 8 bits

Handwritten binary subtraction showing a borrow (b0) and the result in two's complement.

									b0
10	0	0	0	0	1	0	1	0	↑
30	0	0	0	1	1	1	1	0	↑
<hr/>									
-20	1	1	1	0	1	1	0	0	
20	0	0	0	1	0	1	0	0	

Multi-precision subtraction example

Assume $[L1, L0] = A$, $[R1, R0] = B$, produce $[L11, L10] = A - B$:

```
sub      L10, L0, R0    ; Low byte
sub      mp L11, L1, R1  ; High byte
subzz    mp L12          ; Overflow?
```

Difference between signed and unsigned: $A - B$, with $A = 10$, $B = 12$

Overflow detection?

$$A - B = A + (-B)$$

Selection: signed min and max

SMAXC

Maximize with C: **Rdst** gets the maximum of the signed ALU result and a signed operand (**OpC**).

<ALUop> Rdst, OpA, OpB smaxc OpC or

smaxc Rdst, OpA, OpC

SMINC

Minimize with C: **Rdst** gets the minimum of the signed ALU result and a signed operand (**OpC**).

<ALUop> Rdst, OpA, OpB sminc OpC or

sminc Rdst, OpA, OpC

Multi-precision selection

In all cases, start from the MSB, and use:

CMP with all unsigned instructions:

maxc, minc, ~~modmaxc, modminc~~, selectc

CMPSWAP with all signed instructions:

smaxc, sminc


However, in the signed case the MSBs have to be swapped.

Example: unsigned multi-precision selection

To select the max of $A = [L2, L1, L0]$ and $B = [L6, L5, L4]$, both unsigned, storing the result in $[L12, L11, L10]$, do:

```
maxc    L12, L2, L6
maxc    L11, L1, L5 cmp
maxc    L10, L0, L4 cmp
```

msb



Example: signed multi-precision selection

To select the max of $A = [L2, L1, L0]$ and $B = [L6, L5, L4]$, both signed, storing the result in $[L12, L11, L10]$, do:

`move L2, #255` ; sample value: $A = -23$

`move L1, #255`

`move L0, #233`

`move L6, #0` ; sample value: $B = 200$

`move L5, #0`

`move L4, #200`

Handwritten binary representations:

- $A = -23$: $[11111111, 11111111, 11111001]$
- $B = 200$: $[00000000, 00000000, 11001000]$

`smaxc L12, L6, L2` ; note the swapped MSBs

`smaxc L11, L1, L5 cmpswap`

`smaxc L10, L0, L4 cmpswap`

Multi-precision comparison

Compare multi-byte numbers starting from the MSB.

Use **ltc** on the MSB, and then again **ltc** with:

CMP on the other bytes, for *unsigned* comparison.

CMPSWAP on the other bytes, for *signed* comparison.

Also, for signed comparison, the high byte operands must be in the opposite order from the others.

To compare for equality, still use **ltc**, but then use

fbeqlatch fbinv bspush

to push the equality latch with the proper assertion.

Example: unsigned multi-precision comparison

Assume three-byte unsigned numbers:

A = [L22, L21, L20] and B = [L19, L18, L17].

To implement

```
if (A == B) {
```

do:

```
addxz  L22, L22 ltc L19  
addxz  L21, L21 ltc L18 cmp  
addxz  L20, L20 ltc L17 cmp  
fbeqlatch fbinv bspush
```

Example: signed multi-precision comparison

Assume three-byte signed numbers:

A = [L22, L21, L20] and B = [L19, L18, L17].

To implement

```
if (A < B) {  
do:
```


```
addxz  L19, L19 sltc L22  
addxz  L21, L21 sltc L18 cmpswap  
addxz  L20, L20 sltc L17 cmpswap bspush
```

TYPOS

The controller's scratch register: `arrtoscr`

ARRTOSCR “ARRay TO controller SCRatch register” writes the result of the instruction into the controller's scratch register as well as into the destination register.

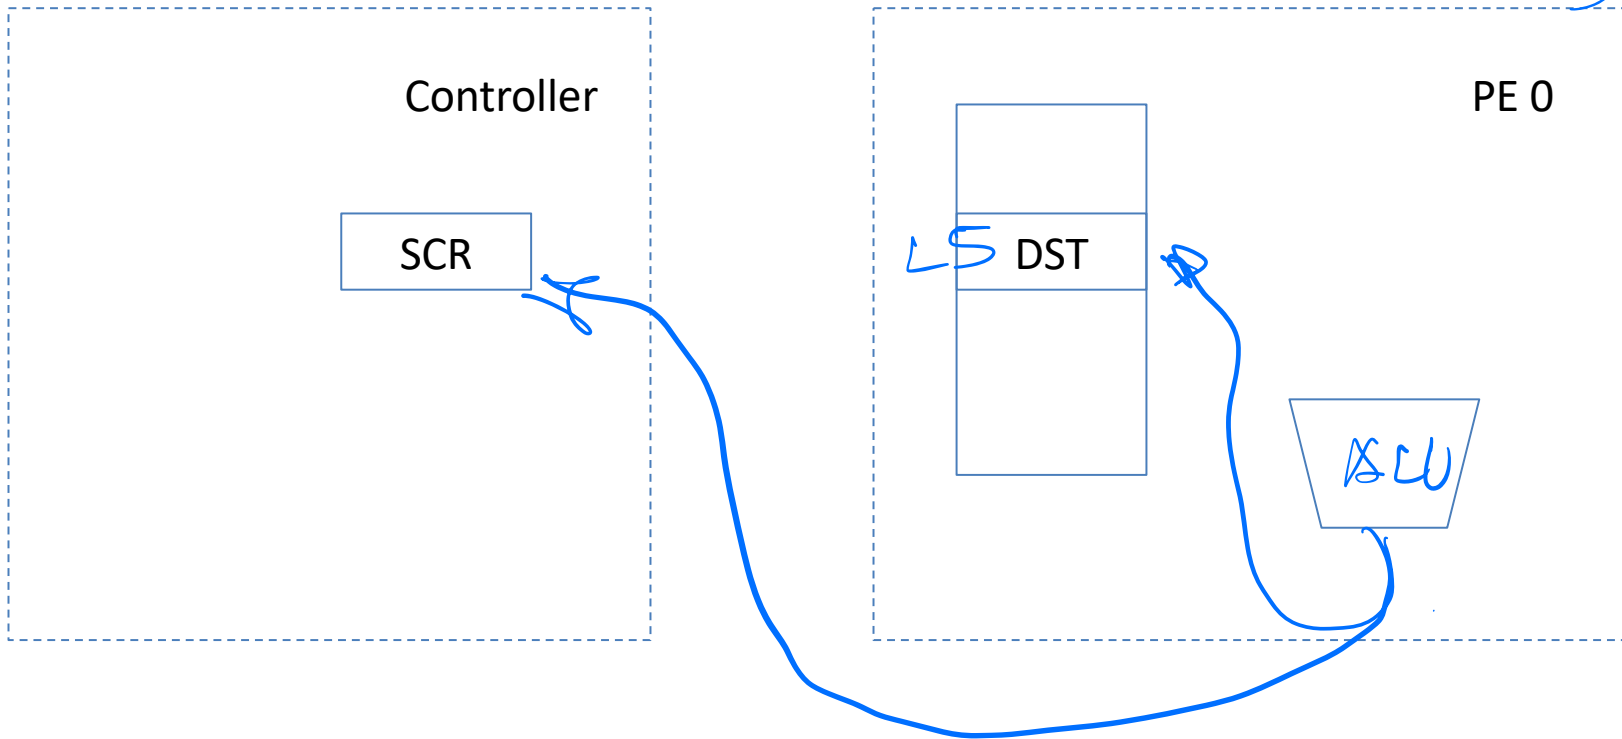
`<instr> Rdst, ... arrtoscr`



- If **Rdst** is a left register, the value will come from the leftmost PE in the array; if **Rdst** is a right register, the value will come from the rightmost PE in the array.
- If the outputting PE is masked, the scratch register (just like the destination register) **SHOULD NOT** be written (see Kasm manual) but a **bug in the simulator makes it write anyway**.

The controller's scratch register: `arrtoscr`

ADD E5...
ARR
TO
SCR



The controller's scratch register: **scrtoarr**

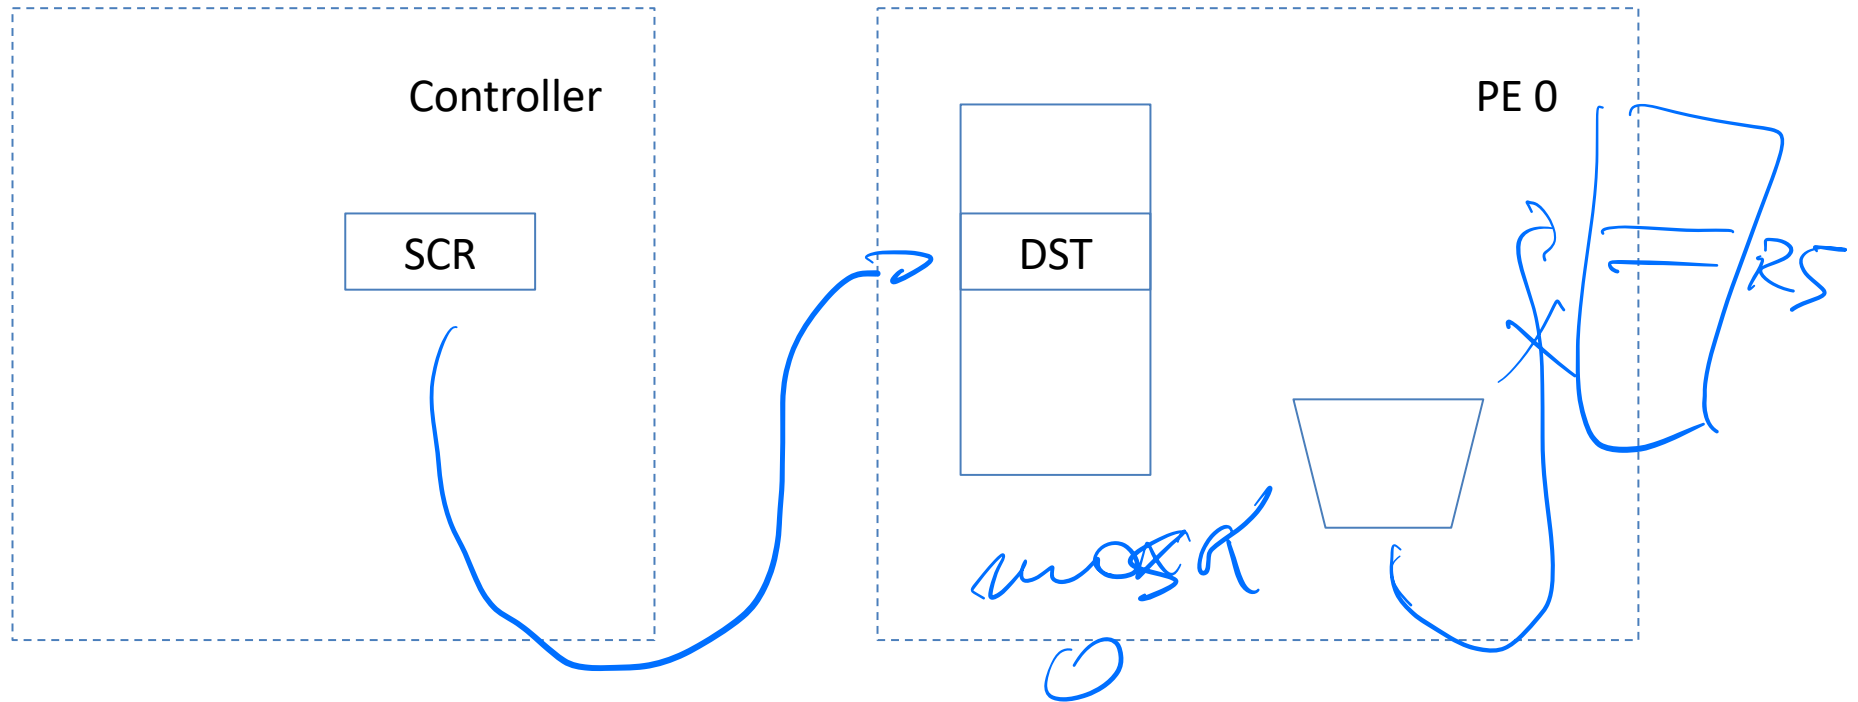
SCRTOARR “controller **SCR**atch register **TO ARR**ay” writes the content of the controller's scratch register into the destination register **Rdst**.

<instr> Rdst, ... scrtoarr


- If **Rdst** is a left register, the value will be written to the right register of the rightmost PE in the array; if **Rdst** is a right register, the value will be written to the left register of the leftmost PE in the array.
- If the receiving PE is masked, the destination register is written anyway.

The controller's scratch register: `scrtoarr`


ADD RS, ..., SCRTOARRN



Loop through scratch



#IMM →
#SCR →



The scratch register: **scrtoimm**, **qtoscr**

SCRTOIMM “Controller **SCR**atch register **TO** array **IMM**ediate”
replaces the array immediate field with the scratch register:

<instr using array imm> scrtoimm

- Can be expressed as **#scr** whenever an array immediate is used.
- The scratch register cannot be written in the instruction immediately before.

QTOSCR “Input **Q**ueue **TO** controller **SCR**atch register” writes a byte from the input queue to the scratch register:

<any array instr> qtoscr

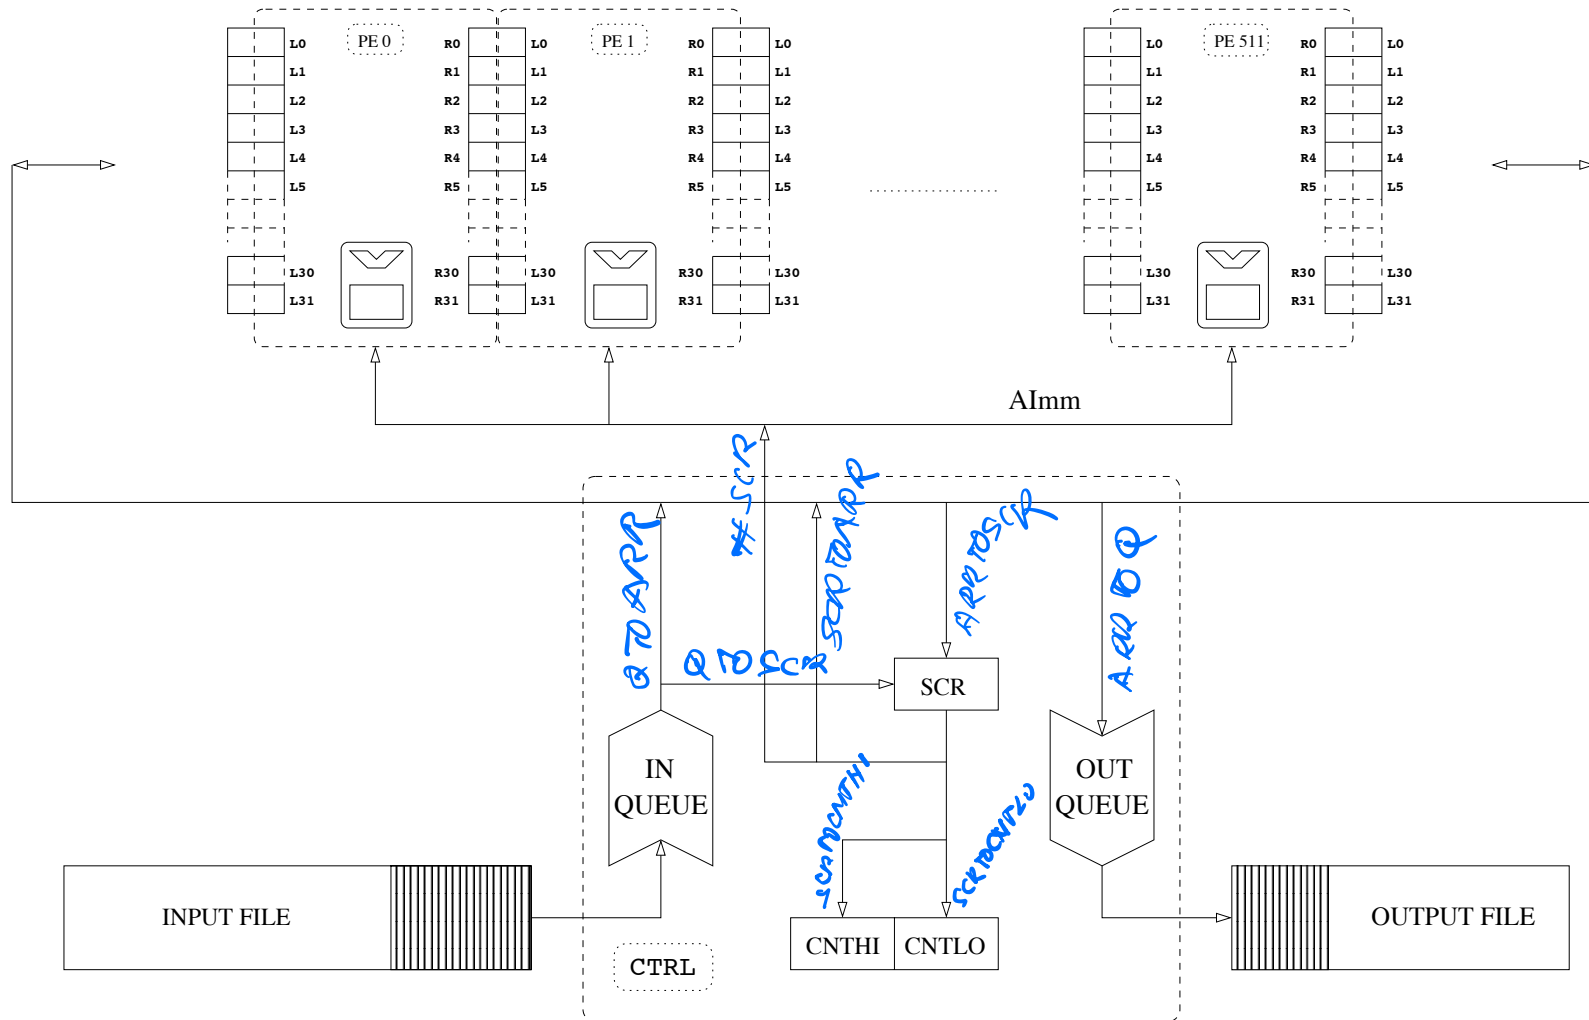
One intervening instruction is needed between **qtoscr** and using the value in the scratch register.

(SCRTOQ) Does not exist! More on this later.

BROADCAST
O(1)
CONSTANT TIME




Controller's scratch register and loop counter



Loop counter stack

The controller's loop counter is a 16-deep *stack* of 16-bit counters:



Setting the loop counter at runtime

Two steps: first set the loop counter, then start the loop. To set the loop counter, there are two ways. The first one is the same as **beginloop**:

CNTPUSH “CouNTER **PUSH**” pushes the controller’s 16-bit immediate onto the counter stack.

<any array instr> **cntpush** CImm *= BEGIN LOOP*

Setting the loop counter at runtime

The second way uses an actual runtime value:

SCRTOCNTLO “SCRatch TO CouNTer LOw byte” writes the scratch register into the low byte of the 16-bit loop counter (top of stack).

<any array instr> scrtocntlo

SCRTOCNTHI “SCRatch TO CouNTer HIgh byte” writes the scratch register into the high byte of the 16-bit loop counter (top of stack).

<any array instr> scrtocnthi



To preserve potential other values already on the stack, one must first make room on the stack by pushing, for instance, a dummy immediate value:

<any array instruction> cntpush 0

Starting the loop

~~FOR~~ = WHILE DO

DO WHILE



To start the loop:

BEGINLOOPSCR Pseudo-instruction that assumes that the counter is already loaded and generates a label to match it with an **endloop** instruction.

<any array instr> beginloopscr

- Due to a design... feature, when loading the counter with a value m , the loop is actually executed $m + 1$ times.
- **beginloopscr** matches with an **endloop** instruction just like **beginloop**
- Can never execute a loop zero times, since it is a do-while loop.

Example of runtime loop

Assume that the number of times we want to loop is a 16-bit number registers [L24, L23] of the leftmost PE.

```
subxz    b1      L23, L23
subxz    mp      L24, L24
move     L23, L23, arrtoscr
```

```
cntpush 0
```

```
nop
```

```
scrtocntlo
```

```
move     L24, L24, arrtoscr
```

```
nop
```

```
nop
```

```
scrtocnthi
```

```
BEGINLOOPSCR
```

```
move     L28, R28, qtoarr ; read row in
```

```
ENDLOOP
```

POD

} 2 INSTR.

} 2 INSTR



```

subxz    b1      L23, L23, arrtoscr
subxz    mp      L24, L24, cntpush 0
nop
scrtocntlo
move      L24, L24, arrtoscr
nop
nop
scrtocnthi

```

```
ENDLOOP    move    L28, R28, qtoarr    ; read row in
```



Nested loops


Loops can be nested to up to 16 levels deep. Make sure to keep that into account when using or nesting macros.

```
BEGINLOOP 10  
  ...  
  BEGINLOOP 20  
    ...  
    BEGINLOOP 30  
      ...  
      ENDLOOP  
    ...  
  ENDLOOP  
  ...  
ENDLOOP
```

Scratch register and I/O operations


The scratch register is always written in an **arrtoq** operation:

move L0, R0, arrtoq



The scratch register can be written directly from the input queue:

<any instr> qtoscr



Scratch register and I/O operations

There is no **scrtoq** instruction. Output from the scratch register has to go through one of the end PEs via the immediate field:

FORCE **move L0, #scr, arrtoq**

or

~~**FORCEQOUT** **move L0, #scr, arrtoq**~~

Local memory

READ Reads a byte from memory into the Memory Data Register (**mdr**).
The byte read is available to the next instruction as **mdr** operand (**OpB**).

Ex:

```
read(L3 + #120)
add      L10, L10, mdr, read(L3 + #121)
add mp   L11, L11, mdr
```

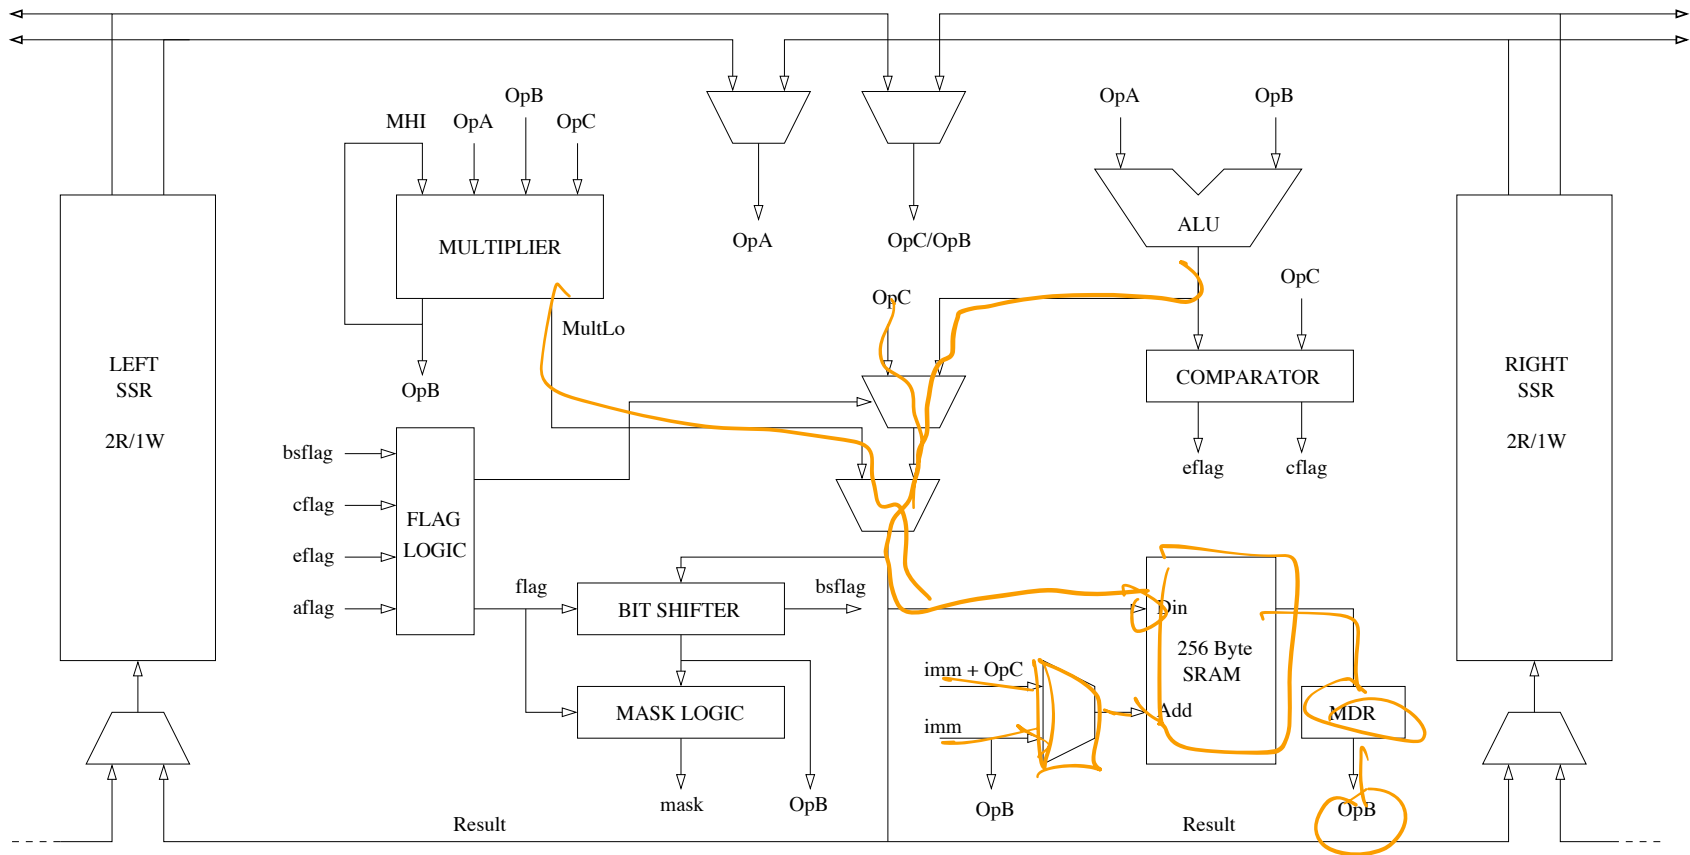


WRITE Writes to memory the result of an operation. Example:


```
move      L10, L10, write(#42)
move      L11, L11, write(#43)
```

Memory reads and writes *cannot* appear in the same instruction.

Kestrel PE diagram



Addressing modes of local memory



Direct

The address is the array 8-bit immediate.

read(AImm)

Indirect

The address is the content of a register.

read(Reg)

Base

The address is the sum of a register and the immediate.

read(Reg + AImm)

The array immediate is always used (so 'indirect' is really 'base'). Therefore even when only a register is specified, no array immediate can be used anyway. NOTE that if you do, the results are undefined, and the assembler will not warn you! Example: **can't do MOVE L0, #3, write(L1)**